



Van Emde Boas Trees

Eduardo Laber

David Sotelo

Lower bounds for heap operations

- Lower bound of $\Omega(n \log n)$ for comparison based sorting.
- As a consequence, at least one among the operations INSERT and EXTRACT-MIN takes $\Omega(\log n)$ time.
- Otherwise HEAPSORT would break down such a lower bound barrier.
- But what about keys in a specific range?

Heaps and sorting in linear time

- There exists faster sorting methods when keys are between **1** and **n** (where **n** is the number of elements).
- In particular, **RADIX SORT** takes $O(n + k)$ time to sort **k** elements from **1** to n^c where **c** is a natural constant.
- van Emde Boas trees support priority queue operations in $O(\log \log u)$ when all keys belong to the range from 0 to $u - 1$.

Van Emde Boas Trees

- Data structure that supports each of the following operations in $O(\log \log u)$ time, where all keys belong to the range $\{0, 1, \dots, u-1\}$:
 - INSERT/DELETE
 - MEMBER
 - MINIMUM/MAXIMUM
 - SUCCESSOR/PREDECESSOR
- We assume, by simplicity, that u is an exact power of two ($u = 2^k$ for some natural number k).
- We also assume that there are not duplicated elements.

Preliminary approaches

In order to gain insight for our problem we shall examine the following preliminary approaches for storing a dynamic set :

- Direct addressing.
- Superimposing a binary tree structure.
- Superimposing a tree of constant height.

Preliminary approaches

In order to gain insight for our problem we shall examine the following preliminary approaches for storing a dynamic set :

- **Direct addressing.**
- Superimposing a binary tree structure.
- Superimposing a tree of constant height.

Direct addressing

- The direct-addressing stores the dynamic set as a bit vector.
- To store a dynamic set of values from the universe $\{0, 1, \dots, u-1\}$ we maintain an array $A[0 .. u-1]$ of u bits.
- Entry $A[x]$ holds **1** if the value x is in the dynamic set and **0** otherwise.
- INSERT, DELETE and MEMBER operations can be performed in $O(1)$ time with this bit vector.
- MINIMUM, MAXIMUM, SUCCESSOR and PREDECESSOR can take $O(u)$ in the worst case since we might have to scan through $O(u)$ elements.

Direct addressing

- $S = \{2, 3, 4, 5, 7, 14, 15\}$ and $u = 16$.

0	0	1	1	1	1	0	1	0	0	0	0	0	0	1	1
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

- $\text{MINIMUM}(A) = 2$
- $\text{MAXIMUM}(A) = 15$
- $\text{SUCCESSOR}(A, 7) = 14$
- $\text{PREDECESSOR}(A, 7) = 5$


Preliminary approaches

In order to gain insight for our problem we shall examine the following preliminary approaches for storing a dynamic set :

- Direct addressing.
- **Superimposing a binary tree structure.**
- Superimposing a tree of constant height.

Superimposing a binary tree structure

- One can short-cut long scans in the bit vector by superimposing a binary tree on the top of it.



Superimposing a binary tree structure


- One can short-cut long scans in the bit vector by superimposing a binary tree on the top of it.
- Entries of the bit vector form the leaves of the binary tree.
- Each internal node contain 1 if and only if in its subtree contains 1.
- Bit stored in an internal node is the logical-or of its children.

Superimposing a binary tree structure

- MINIMUM:
 - **Algorithm:** start at the root and head down toward the leaves, always taking the **leftmost** node containing 1.
 - **Time complexity:** $O(\log u)$
- MAXIMUM:
 - **Algorithm:** start at the root and head down toward the leaves, always taking the **rightmost** node containing 1.
 - **Time complexity:** $O(\log u)$


Superimposing a binary tree structure

- MINIMUM operation:




Superimposing a binary tree structure

- MINIMUM operation:




Superimposing a binary tree structure

- MINIMUM operation:




Superimposing a binary tree structure

- MINIMUM operation:




Superimposing a binary tree structure

- MINIMUM operation:



Superimposing a binary tree structure


- MINIMUM operation:



Minimum found!


Superimposing a binary tree structure

- MAXIMUM operation:




Superimposing a binary tree structure

- MAXIMUM operation:




Superimposing a binary tree structure

- MAXIMUM operation:




Superimposing a binary tree structure

- MAXIMUM operation:




Superimposing a binary tree structure

- MAXIMUM operation:



Superimposing a binary tree structure

- MAXIMUM operation:




Maximum found!

Superimposing a binary tree structure

- PREDECESSOR (A, x) :
 - Start at the leaf indexed by **x** and head up toward the root until we enter a node from the **right** and this node has a 1 in its **left** child **z**.
 - Head down through node **z**, always taking the **rightmost** node containing a 1.
 - **Time complexity:** $O(\log u)$.
- SUCCESSOR (A, x) :
 - Start at the leaf indexed by **x** and head up toward the root until we enter a node from the **left** and this node has a 1 in its **right** child **z**.
 - Head down through node **z**, always taking the **leftmost** node containing a 1.
 - **Time complexity:** $O(\log u)$.


Superimposing a binary tree structure

- PREDECESSOR (A, 7):




Superimposing a binary tree structure

- PREDECESSOR (A, 7):




Superimposing a binary tree structure

- PREDECESSOR (A, 7):




Superimposing a binary tree structure

- PREDECESSOR (A, 7):




Superimposing a binary tree structure

- PREDECESSOR (A, 7):



Superimposing a binary tree structure


- PREDECESSOR (A, 7):



PREDECESSOR (A, 7) = 5


Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):




Superimposing a binary tree structure

- SUCCESSOR (A, 7):



Superimposing a binary tree structure

- SUCCESSOR (A, 7):




SUCCESSOR (A, 7) = 14

Superimposing a binary tree structure

- INSERTION (A, x) :
 - Start at the leaf indexed by x and head up toward the root storing 1 in all nodes belonging to the simple path from x to the root.
 - **Time complexity:** $O(\log u)$.
- DELETION (A, x) :
 - Start at the leaf indexed by x and head up toward the root, recomputing the bit in each internal node on the simple path from x to the root as the logical-or of its two children.
 - **Time complexity:** $O(\log u)$.


Superimposing a binary tree structure

- INSERTION (A, 9):




Superimposing a binary tree structure

- INSERTION (A, 9):




Superimposing a binary tree structure

- INSERTION (A, 9):




Superimposing a binary tree structure

- INSERTION (A, 9):




Superimposing a binary tree structure

- INSERTION (A, 9):



Superimposing a binary tree structure


- INSERTION (A, 9):



INSERTION COMPLETED!


Superimposing a binary tree structure

- DELETION (A, 7):




Superimposing a binary tree structure

- DELETION (A, 7):




Superimposing a binary tree structure

- **DELETION (A, 7):**




Superimposing a binary tree structure

- **DELETION (A, 7):**




Superimposing a binary tree structure

- DELETION (A, 7):



Superimposing a binary tree structure

- DELETION (A, 7):



DELETION COMPLETED!

Superimposing a binary tree structure

- This approach is only marginally better than just using a balanced search tree.
- MEMBER operation can be performed in $O(1)$ time, what is better than $O(\log n)$ for BST.
- If the number of elements n is much smaller than size of the universe u than BST would be faster for all the other operations.


Preliminary approaches

In order to gain insight for our problem we shall examine the following preliminary approaches for storing a dynamic set :

- Direct addressing.
- Superimposing a binary tree structure.
- **Superimposing a tree of constant height.**


Superimposing a tree of constant height

- Let us assume that $u = 2^{2^k}$ for some natural k .
- Consider that we superimpose a tree of degree $u^{1/2}$
- The height of the resulting tree would be always 2.




Superimposing a tree of constant height

- Again, each internal node stores the logical-or of the bits within its subtree.
- The $u^{1/2}$ internal nodes at depth 1 summarize each group (logical-or) of $u^{1/2}$ values.
- These internal nodes can be seen as an array **summary[0 .. $u^{1/2} - 1$]**



Superimposing a tree of constant height

- Given an index x , one can find the cluster that x belongs by the expression $\lfloor x / \sqrt{u} \rfloor$
- Operations INSERT and MEMBER can be performed in $O(1)$ time.
- We can use **summary** array to perform all other operations MAX/MIN/PRED/SUC/DEL in $O(u^{1/2})$ time.



Exercise 1

- Modify the data structures introduced so far in order to support duplicate keys.

Exercise 1

- Modify the data structures introduced so far in order to support duplicate keys.
- **Solution:** keep up one integer in the leaves instead of a single bit for counting repetitions. Internal nodes are not affected by this change.

Exercise 2

- Suppose that we superimpose a tree of degree $u^{1/k}$, instead of $u^{1/2}$, where $k > 1$ is an integer constant. What would be the height of such a tree and how long would each of the operations take?

Exercise 2

- Suppose that we superimpose a tree of degree $u^{1/k}$, instead of $u^{1/2}$, where $k > 1$ is an integer constant. What would be the height of such a tree and how long would each of the operations take?
- **Solution:** The height of the tree would be k and each operation would take $k \cdot u^{1/k}$

An interesting recurrence

- Since we are interested on achieving running times of **$O(\log \log u)$** let us think about how to obtain these bounds.
- The following recurrence will guide our search for a data structure:

$$T(u) = T(\sqrt{u}) + O(1)$$

An interesting recurrence

$$T(u) = T(\sqrt{u}) + O(1)$$

$$m = \lg u \Rightarrow 2^m = u$$

$$\Rightarrow T(u) = T(2^m)$$

$$\Rightarrow T(2^m) = T(2^{m/2}) + O(1)$$

$$S(m) = T(2^m)$$

$$\Rightarrow S(m) = S(m/2) + O(1)$$


$$\Rightarrow S(m) = \lg m$$

$$\Rightarrow T(u) = T(2^m) = S(m) = \lg m = \lg \lg u$$

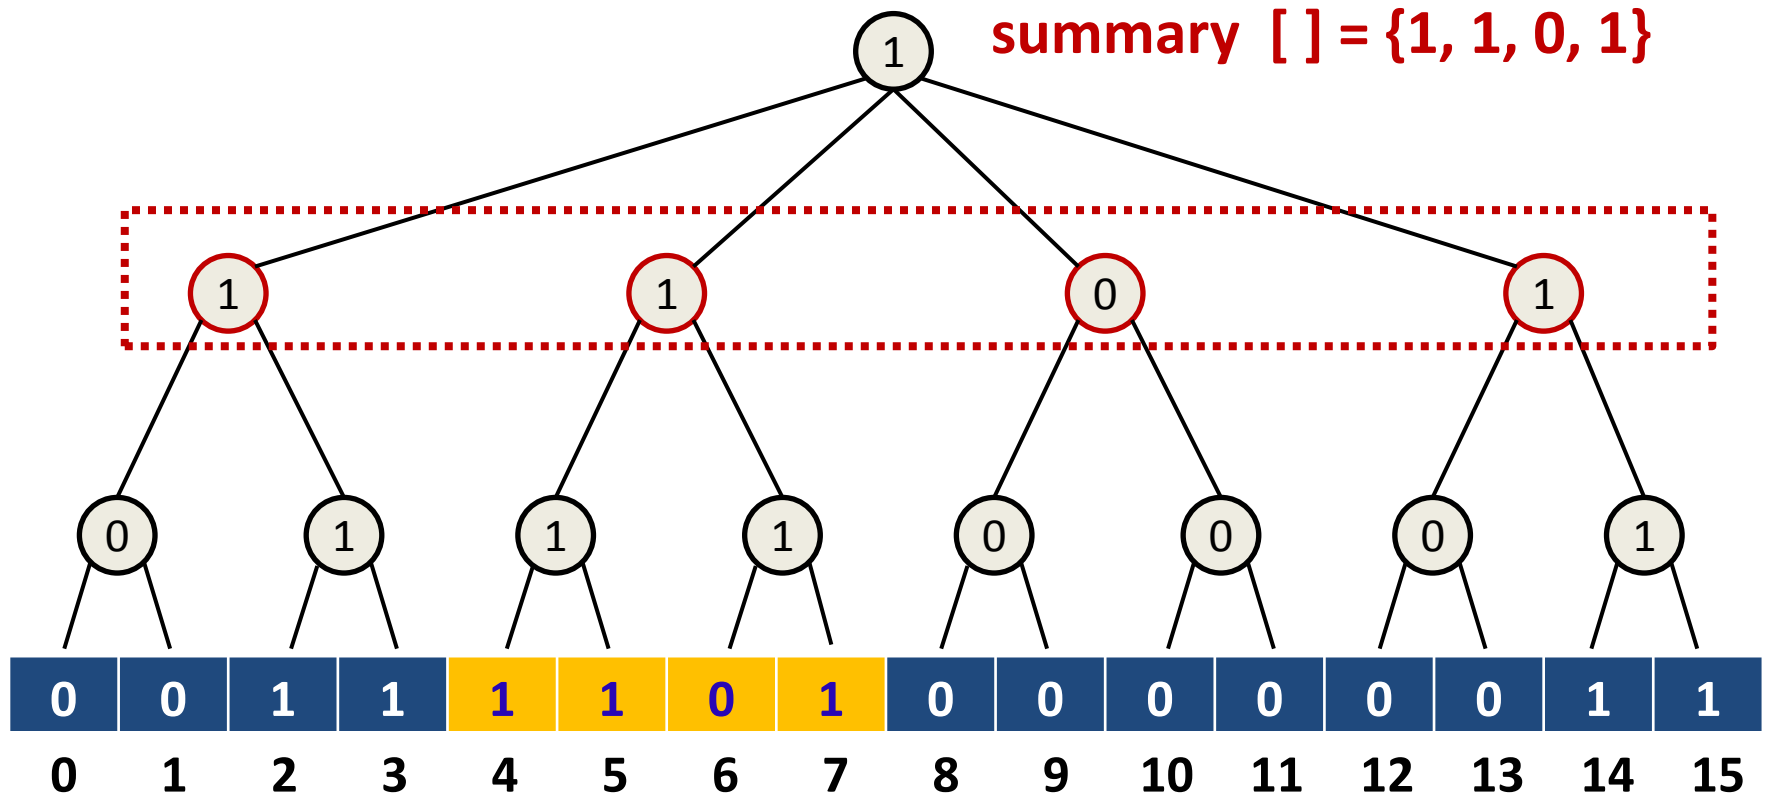
Prototype of van Emde Boas

- The main idea is superimposing a tree of variable degree $u^{1/k}$.
- Starting with an universe of size u we make structures of $u^{1/2}$ items, which themselves contain structures of $u^{1/4}$ items, which hold structures of $u^{1/8}$ items...
- This subdivision process stops when there are only 2 items.
- For simplicity we assume that $u = 2^{2^k}$ for some integer k .

Prototype of van Emde Boas



Prototype of van Emde Boas

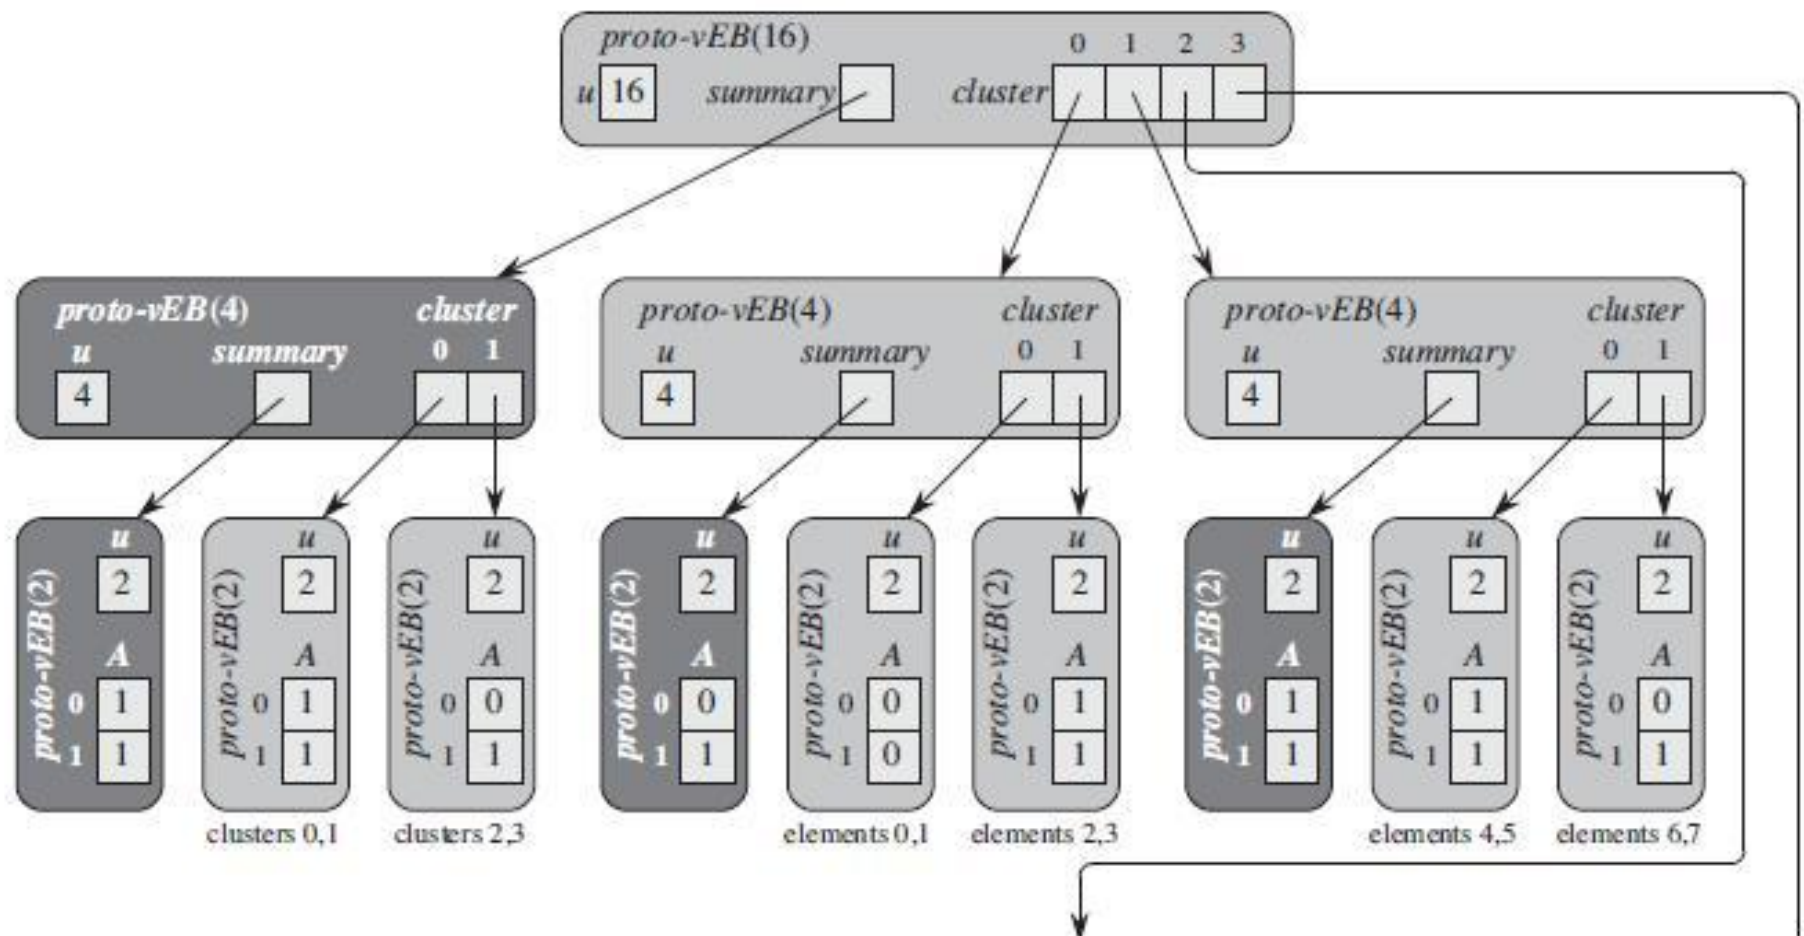


Prototype of van Emde Boas

Let $u = 2^{2^k}$ be the size of the universe for some integer $k \geq 0$.

- If $u = 2$ then it is the basis size, and the **proto-vEB(u)** contains the array $A[0..1]$ of two bits.
- Otherwise the **proto-vEB(u)** structure is defined recursively as follows:
 - It contains a pointer named **summary** to a **proto-vEB($u^{1/2}$)** structure.
 - Its contains an array **cluster**[0 .. $u^{1/2} - 1$] of $u^{1/2}$ pointers, each to a **proto-vEB($u^{1/2}$)** structure.

Prototype of van Emde Boas



Clusters and positions

- A given value x resides in **cluster number** $\lfloor x / \sqrt{u} \rfloor$
- Since that x is a binary integer of $\lg u$ bits, its cluster number is given by the **most significant** $(\lg u)/2$ bits of x .
- The **position** in which x appears in its cluster is given by the expression $x \bmod \sqrt{u}$
- As a consequence, the position of x inside its cluster is given by the **least significant** $(\lg u)/2$ bits of x

Clusters and positions

The following functions will be necessary to explore the **proto-vEB** structure:

- $\text{high}(x) = \lfloor x / \sqrt{u} \rfloor$
- $\text{low}(x) = x \bmod \sqrt{u}$
- $\text{index}(x, y) = x\sqrt{u} + y$

Member operation on proto-vEB

PROTO-vEB-MEMBER (V, x)

```
1   if V.u == 2
2       return V.A[x]
3   else
4       return PROTO-vEB-MEMBER (V.cluster[high(x)], low(x))
```

Member operation on proto-vEB

- Running time:

$$T(u) = T(\sqrt{u}) + O(1)$$
$$\Rightarrow T(u) = O(\lg \lg u)$$

Finding the minimum of proto-vEB

PROTO-vEB-MINIMUM (V)

```
1  if V.u == 2
2      if V.A[0] == 1
3          return 0
4      elseif V.A[1] == 1
5          return 1
6      else return NIL
7  else min-cluster = PROTO-vEB-MINIMUM (V.summary)
8      if min-cluster == NIL
9          return NIL
10     else offset = PROTO-vEB-MINIMUM (V.cluster[min-cluster])
11     return index(min-cluster, offset)
```

Finding the minimum of proto-vEB

- Running time:

$$T(u) = 2T(\sqrt{u}) + O(1)$$

$$\Rightarrow T(2^m) = 2T(2^{m/2}) + O(1)$$

$$\Rightarrow S(m) = 2S(m/2) + O(1)$$

$$\Rightarrow S(m) = \Theta(m)$$

$$\Rightarrow T(u) = T(2^m) = S(m) = \Theta(m) = \Theta(\lg u)$$

Successor operation on proto-vEB

PROTO-vEB-SUCCESSOR (V, x)

```
1  if V.u == 2
2      |   if x == 0 and V.A[1] == 1
3          |   return 1
4      |   else return NIL
5  else offset = PROTO-vEB-SUCCESSOR (V.cluster[high(x)], low(x))
6      |   if offset ≠ NIL
7          |   return index (high(x), offset)
8      |   else
9          |   succ-cluster = PROTO-vEB-SUCCESSOR (V.summary, high(x))
10         |   if succ-cluster == NIL
11             |   return NIL
12         |   else offset = PROTO-vEB-MINIMUM (V.cluster[succ-cluster])
13         |   return index(succ-cluster, offset)
```

Successor operation on proto-vEB

- Running time:

$$T(u) = 2T(\sqrt{u}) + O(\lg \sqrt{u})$$

$$\Rightarrow T(2^m) = 2T(2^{m/2}) + O(m)$$

$$\Rightarrow S(m) = 2S(m/2) + O(m)$$

$$\Rightarrow S(m) = \Theta(m \lg m)$$

$$\Rightarrow T(u) = T(2^m) = S(m)$$

$$\Rightarrow T(u) = \Theta(m \lg m) = \Theta(\lg u \lg \lg u)$$

Insertion on proto-vEB

PROTO-vEB-INSERT (V, x)

- 1 **if** $V.u == 2$
- 2 $V.A[x] = 1$
- 3 **else**
- 4 PROTO-vEB-INSERT ($V.cluster[high(x)], low(x)$)
- 5 PROTO-vEB-INSERT ($V.summary, high(x)$)

Insertion on proto-vEB

- Running time:

$$T(u) = 2T(\sqrt{u}) + O(1)$$

$$\Rightarrow T(2^m) = 2T(2^{m/2}) + O(1)$$

$$\Rightarrow S(m) = 2S(m/2) + O(1)$$

$$\Rightarrow S(m) = \Theta(m)$$

$$\Rightarrow T(u) = T(2^m) = S(m) = \Theta(m) = \Theta(\lg u)$$

The “REAL” van Emde Boas Tree

- Proto-vEB structure was close to what is needed to achieve $O(\lg \lg n)$ running times.
- It requires too many (2 or more) recursions in most of the operations (time complexity bounds allow only one recursion).
- The “REAL” van Emde Boas tree stores some extra information in order to make at most one recursive call on each of the operations.

The universe size $u = 2^k$

- The universe size $u = 2^{2^k}$ in **proto-vEB** was very restrictive :
 - If $k = 1$ then $u = 4$
 - If $k = 2$ then $u = 16$
 - If $k = 3$ then $u = 256$
 - If $k = 4$ then $u = 65536$
 - If $k = 5$ then $u = 4294967296$
 - If $k = 6$ then $u = 18446744073709551616$
- In the **proto-vEB**, any $u \neq 2^{2^k}$ must be transformed into an u' such that $u' = \Omega(u^2)$
- The **real-vEB** allow the universe size u to be **any exact power of two**. In other words, $u = 2^k$ for an integer $k > 0$

The universe size $u = 2^k$

- Since that the universe size is any exact power of two, $u^{1/2}$ is not necessarily an integer.
- However, this change affects the calculation of functions **high**, **low** and **index** that uses most/least significant bits.
 - $\text{high}(x) = \lfloor x / \sqrt{u} \rfloor$
 - $\text{low}(x) = x \bmod \sqrt{u}$
 - $\text{index}(x, y) = x\sqrt{u} + y$

The universe size $u = 2^k$

- For convenience, we introduce the concepts of **lower square root** and **upper square root**:
 - **Lower square root:** $\downarrow \sqrt{u} = 2^{\lfloor (\lg u)/2 \rfloor}$
 - **Upper square root:** $\uparrow \sqrt{u} = 2^{\lceil (\lg u)/2 \rceil}$
- **Examples:**
 - $\downarrow \sqrt{32} = 2^{\lfloor (\lg 32)/2 \rfloor} = 2^2 = 4$
 - $\uparrow \sqrt{32} = 2^{\lceil (\lg 32)/2 \rceil} = 2^3 = 8$

Redefining high, low and index

- Since that the universe size is any exact power of two, $u^{1/2}$ is not necessarily an integer.
- However, this change affects the calculation of functions **high**, **low** and **index** that uses most/least significant bits.
 - $\text{high}(x) = \lfloor x / \downarrow \sqrt{u} \rfloor$
 - $\text{low}(x) = x \bmod \downarrow \sqrt{u}$
 - $\text{index}(x) = x \downarrow \sqrt{u} + y$

Redefining high, low and index

■ Example


If $u = 36$ and $x = 23$, then:

- $\text{high}(23) = \lfloor 23 / \downarrow \sqrt{32} \rfloor = \lfloor 23 / 4 \rfloor = 5$
- $\text{low}(23) = 23 \bmod \downarrow \sqrt{32} = 23 \bmod 4 = 3$
- $\text{index}(5, 3) = 5 \downarrow \sqrt{32} + 3 = 5 \cdot 4 + 3 = 23$

The “REAL” van Emde Boas Tree

- The **real-vEB** modifies the proto-VEB, containing two new attributes:
 - **min** stores the minimum element in the real-vEB.
 - **max** stores the maximum element in the real-vEB.
- The attribute **summary** points to a **real-vEB**($\uparrow u^{1/2}$) tree.
- The array **cluster** points to $\uparrow u^{1/2}$ **real-vEB**($\downarrow u^{1/2}$) trees.
- **Important:** The element stored in **min** does not appear in any of the recursive real-vEB trees that the *cluster* array points to.

The “REAL” van Emde Boas Tree

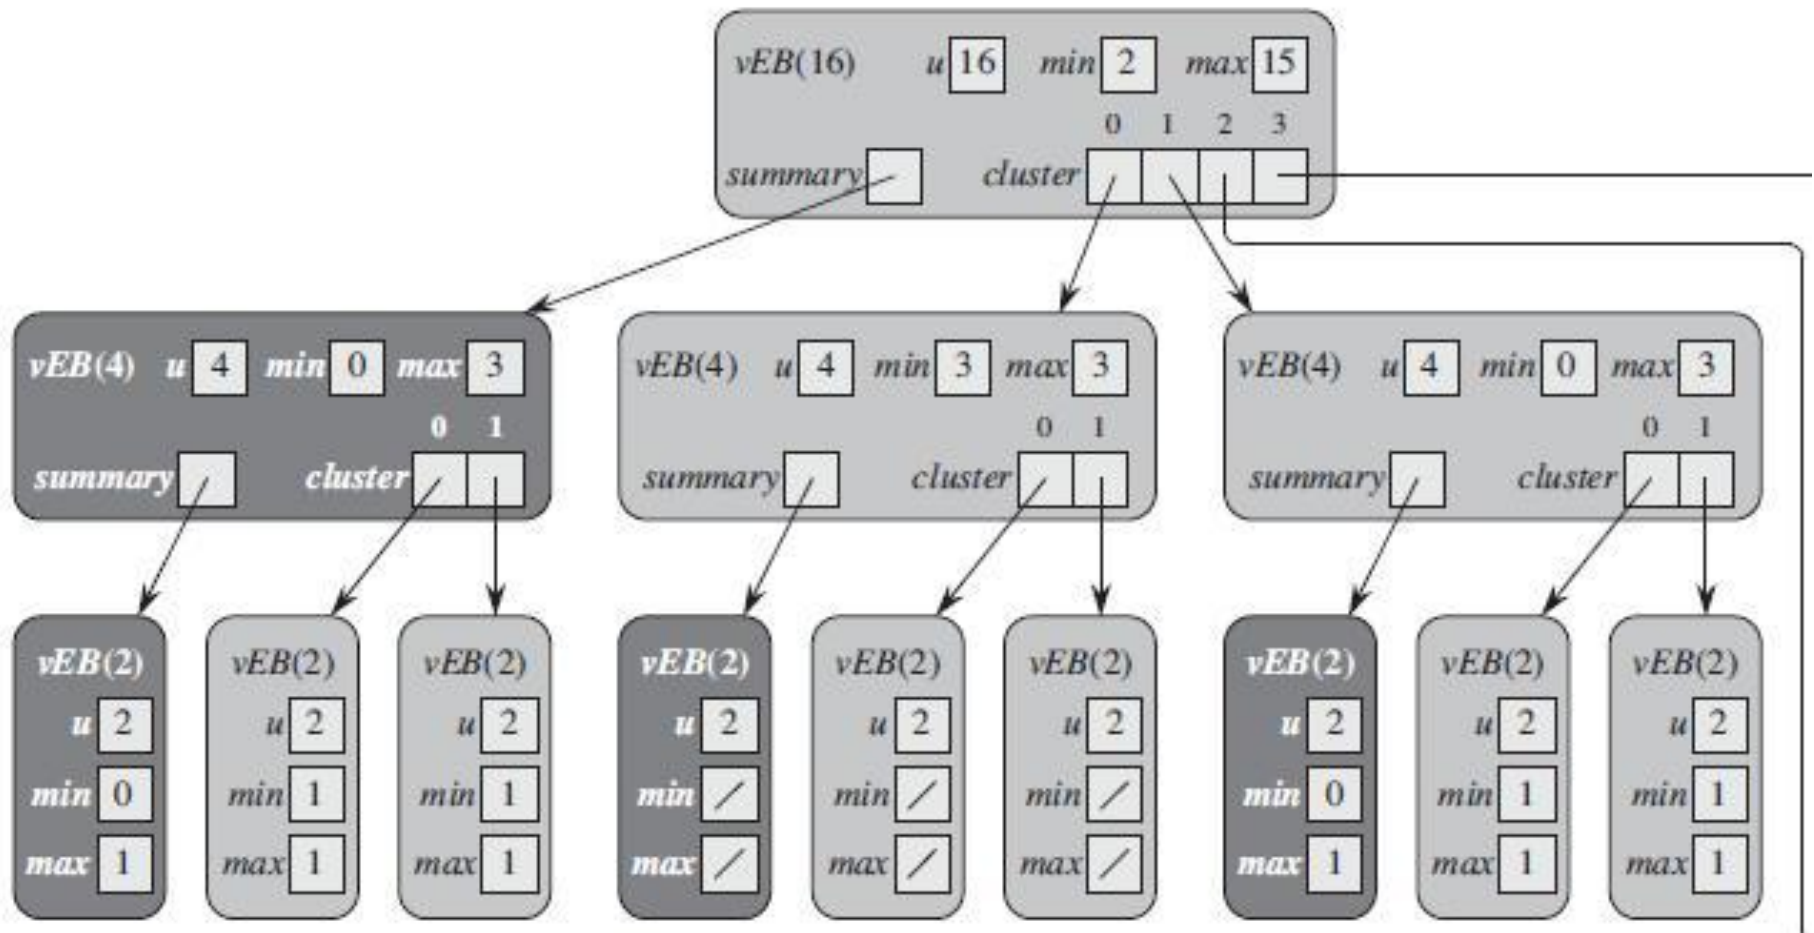


The “REAL” van Emde Boas Tree

- The elements stored in a **real-vEB(u)** tree **V.min** plus all the elements stored in the **real-vEB($\downarrow u^{1/2}$)** trees pointed to by **V.cluster[0 .. ($\uparrow u^{1/2}$) - 1]**
- When a real-vEB contains two or more elements the element stored in **min** does not appear in any of the clusters but the stored in **max** does (unless it is equal to **min**).
- Since the basis size is 2, real-vEB(2) does not need the array **A** of the corresponding proto-vEB(2). The attributes **min** and **max** can be used for it.
- **Important:** In the real-vEB, all the elements can be determined from the **min** and **max** attributes.

The “REAL” van Emde Boas Tree

- **Example:** real-vEB(16) holding the set {2, 3, 4, 5, 7, 14, 15}



Minimum and Maximum on vEB tree

vEB-TREE-MINIMUM (V, x)

1 return V.min

vEB-TREE-MAXIMUM (V, x)

1 return V.max

- Running time: $O(1)$

Member operation on vEB tree

vEB-TREE-MEMBER (V, x)

```
1  if x == V.min or x == V.max
2  |   return TRUE
3  elseif V.u == 2
4  |   return FALSE
5  else return vEB-TREE-MEMBER (V.cluster[high(x)], low(x))
```

- Running time:

$$T(u) = T(\sqrt{u}) + O(1)$$

$$\Rightarrow T(u) = O(\lg \lg u)$$

Successor operation on proto-vEB

PROTO-vEB-SUCCESSOR (V, x)

```
1  if V.u == 2
2      if x == 0 and V.A[1] == 1
3          return 1
4      else return NIL
5  else offset = PROTO-vEB-SUCCESSOR (V.cluster[high(x)], low(x))
6      if offset ≠ NIL
7          return index (high(x), offset)
8      else
9          succ-cluster = PROTO-vEB-SUCCESSOR (V.summary, high(x))
10         if succ-cluster == NIL
11             return NIL
12         else offset = PROTO-vEB-MINIMUM (V.cluster[succ-cluster])
13             return index(succ-cluster, offset)
```

Successor operation on vEB tree

vEB-TREE-SUCCESSOR (V, x)

```
1  if V.u == 2
2    |   if x == 0 and V.max == 1
3    |       return 1
4    |   else return NIL
5  elseif V.min ≠ NIL and x < V.min
6    |   return V.min
7  else max-low = vEB-TREE-MAXIMUM (V.cluster[high(x)])
8    |   if max-low ≠ NIL and low(x) < max-low
9    |       offset = vEB-TREE-SUCCESSOR(V.cluster[high(x)], low(x))
10   |   return index (high(x), offset)
11  else succ-cluster = vEB-TREE-SUCCESSOR(V.summary, high(x))
12   |   if succ-cluster == NIL
13   |       return NIL
14   |   else offset = vEB-TREE-MINIMUM (V.cluster[succ-cluster])
15   |       return index(succ-cluster, offset)
```

Successor operation on vEB tree

- Running time:

$$T(u) \leq \max \left\{ T(\downarrow \sqrt{u}), T(\uparrow \sqrt{u}) \right\} + O(1)$$

$$\Rightarrow T(u) \leq T(\uparrow \sqrt{u}) + O(1)$$

$$m = \lg u$$

$$\Rightarrow T(2^m) \leq T(2^{\lceil m/2 \rceil}) + O(1)$$

$$(\forall m \geq 2) \lceil m/2 \rceil \leq 2m/3$$

$$\Rightarrow T(2^m) \leq T(2^{2m/3}) + O(1)$$

Successor operation on vEB tree

- Running time:

$$T(2^m) \leq T(2^{2m/3}) + O(1)$$

$$S(m) = T(2^m)$$

$$\Rightarrow S(m) \leq S(2m/3) + O(1)$$

$$\Rightarrow S(m) = O(\lg m)$$

$$\Rightarrow T(u) = T(2^m) = S(m) = O(\lg m) = O(\lg \lg u)$$

Predecessor operation on vEB tree

vEB-TREE-PREDECESSOR (V, x)

```
1  if V.u == 2
2  |   if x == 1 and V.min == 0
3  |   |   return 0
4  |   else return NIL
5  elseif V.max ≠ NIL and x > V.max
6  |   return V.max
7  else min-low = vEB-TREE-MINIMUM (V.cluster[high(x)])
8  |   if min-low ≠ NIL and low(x) > min-low
9  |   |   offset = vEB-TREE-PREDECESSOR(V.cluster[high(x)], low(x))
10 |   |   return index (high(x), offset)
11 else pred-cluster = vEB-TREE-PREDECESSOR(V.summary, high(x))
12 |   if pred-cluster == NIL
13 |   |   if V.min ≠ NIL and x > V.min
14 |   |   |   return V.min
15 |   |   else return NIL
16 |   else offset = vEB-TREE-MAXIMUM (V.cluster[pred-cluster])
17 |   |   return index(succ-cluster, offset)
```

Predecessor operation on vEB tree

- Running time:

$$T(u) \leq \max \left\{ T\left(\downarrow \sqrt{u}\right), T\left(\uparrow \sqrt{u}\right) \right\} + O(1)$$

$$\Rightarrow T(u) = O(\lg \lg u)$$

Insert operation on proto-vEB

PROTO-vEB-INSERT (V, x)

- 1** **if** $V.u == 2$
- 2** $V.A[x] = 1$
- 3** **else**
- 4** **PROTO-vEB-INSERT** ($V.cluster[high(x)], low(x)$)
- 5** **PROTO-vEB-INSERT** ($V.summary, high(x)$)

Insert operation on vEB tree

- The insertion operation can be split into two cases.
The first one occurs when the cluster is empty.

vEB-EMPTY-TREE-INSERT (V, x)

1 $V.\text{min} = x$

2 $V.\text{max} = x$

- In the second case, one must avoid the two recursive calls, one for the summary and another for a cluster. How can we do it?
- **Answer:** check if the corresponding cluster is empty or not.

Insert operation on vEB tree

vEB-TREE-INSERT (V, x)

```
1  if V.min == NIL
2    vEB-EMPTY-TREE-INSERT (V, x)
3  else
4    if x < V.min
5      exchange x with V.min
6    if V.u > 2
7      if vEB-TREE-MINIMUM (V.cluster[high(x)]) == NIL
8        vEB-TREE-INSERT(V.summary, high(x))
9        vEB-EMPTY-TREE-INSERT (V.cluster[high(x)], low(x))
10     else vEB-TREE-INSERT (V.cluster[high(x)], low(x))
11  if x > V.max
12    V.max = x
```

Insert operation on vEB tree

- Running time:

$$T(u) \leq \max \left\{ T\left(\downarrow \sqrt{u}\right), T\left(\uparrow \sqrt{u}\right) \right\} + O(1)$$
$$\Rightarrow T(u) = O(\lg \lg u)$$